

Chapter 7 n-step Bootstrapping

🕒 Created	@December 4, 2023 9:08 AM
🏷️ Tags	

written by Yee

7-1. n-step TD Prediction

n-step TD 大概為 MC 跟 TD 的綜合體，MC更新的時候是在整個episode跑完之後才去做更新，而TD只會基於上一個值的變化做更新。我們若是可以基於多個step，但又不想要跑完整個episode在做更新，比如2-step、3-step等等多-step的更新。這些n-step更新仍然適用TD因為她仍然基於先前幾步的差值去做更新。有別於 $TD(0)$ 的每步更新，現在更新得要延遲 n-step。或是我們可以稱之為 **n-step TD method**。

假設有一個episode 有state-reward sequence, $S_t, R_{t+1}, S_{t+1}, R_{t+2}, S_{t+2}, \dots, R_T, S_T$ ，我們知道MC是需要到episode結束才能更新資料。

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T$$

有別於MC目標是return，one-step return 是的目標是reward及下個state的期望值。

$$G_{t:t+1} = R_{t+1} + \gamma V_t(S_{t+1})$$

這個 $t : t + 1$ 的寫法表示了 擷取 time t 到 $t + 1$ 的 reward總值的 return，之後再以 $\gamma V_t(S_{t+1})$ 取代 $\gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T$ ，我們以此類推到下一個step，以下為**two-step return**

$$G_{t:t+2} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$$

接下來再繼續推到 **n-step return**

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n}) \quad (7.1)$$

for all n, t such that $n \geq 1$ and $0 \leq t < T - n$

如果 $t + n \geq T$ ，那所有缺少的元素都補 0，這時候 n-step return 定義回原來的 G_t Full return

$$(G_{t:t+n} = G_t \text{ if } t + n \geq T)$$

n-step如果 $n > 1$ 會需要未來的Reward，需要直到 R_{t+n} 之後並計算 V_{t+n-1} ，再 $t + n$ 的時間才能得到下列這兩個的值。

$$V_{t+n}(S_t) = V_{t+n-1}(S_t) + \alpha[G_{t:t+n} - V_{t+n-1}(S_t)] \quad (7.2)$$

我們稱之為 **n-step TD**。在一開始的 $n - 1$ step 我們不會改動任何東西，再最後到terminal state後在一次更新之後還沒更新過的地方。

n-step TD for estimating $V \approx V_\pi$

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can take their index mod $n + 1$

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$
 $T \leftarrow \infty$
 Loop for $t = 0, 1, 2, \dots$:
 If $t < T$, then:
 Take an action according to $\pi(\cdot|S_t)$
 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}
 If S_{t+1} is terminal, then $T \leftarrow t + 1$
 $\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being update)
 If $\tau \geq 0$:
 $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$
 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$
 $V(S_\tau) \leftarrow V(S_\tau) + \alpha[G - V(S_\tau)]$
 Until $\tau = T - 1$

7.2 n-step Sarsa (on-policy)

這個 method 僅僅是把 prediction 改成 control method，然後使用的是 ϵ -greedy policy，Sarsa 是在 action 下結束而非 state。我們重新定義 n-step Sarsa 的 預估action values：

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n}) \quad (7.4)$$

$$n \geq 1, 0 \leq t < T - n$$

$$\text{with } G_{t:t+n} = G_t \text{ if } t + n \geq T$$

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha[G_{t:t+n} - Q_{t+n-1}(S_t, A_t)] \quad (7.5)$$

$$0 \leq t < T$$

這個就稱作 n-step Sarsa

n-step Sarsa for estimating $Q \approx q_\pi$ or q_*

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be ϵ -greedy w.r.t Q , or to a fixed given policy

Algorithm parameter: step size $\alpha \in (0, 1]$, small $\epsilon > 0$, a positive integer n

All store and access operation (for S_t, A_t, R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq \text{terminal}$
 Select and store an action $A_0 \sim \pi(\cdot|S_0)$
 $T \leftarrow \infty$
 Loop for $t = 0, 1, 2, \dots$:
 If $t < T$, then:

Take action A_t
 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}
 If S_{t+1} is terminal
 then $T \leftarrow t + 1$
 else:
 Select and store an action $A_{t+1} \sim \pi(\cdot|S_{t+1})$
 $\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being update)
 If $\tau \geq 0$:
 $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$
 If $r + n < T$, then: $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$
 $Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha[G - Q(S_\tau, A_\tau)]$
 If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is ϵ -greedy w.r.t Q
 Until $\tau = T - 1$

那 Expected Sarsa呢?

在前面behavior的過程仍舊跟Sarsa一樣是線性的，除了最後一個元素需要分支計算下一個state的所有期望值總和。 \bar{V}_t

$$G_{t:t+n} = R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \bar{V}_{t+n-1}(S_{t+n}) \quad (7.9)$$

$t + n < T$

($G_{t:t+n} = G_t$ if $t + n \geq T$) \bar{V}_t 是 state s 的期望預估值 (expected approximate value)

7.3 n-step Off-policy Sarsa

回顧 off-policy 是讓 policy π 從 behavior policy b 上學習。 π 通常是以目前 Action-value function所組成的 greedy policy， b 則通常是比較具探索性(exploratory) 的 policy，通常是 ϵ -greedy policy。為了可以使用 b 產出的資料，我們需要計算走這步的相對機率(Section 5.5)。回顧 importance sampling ratio

$$V_{t+n}(S_t) = V_{t+n-1}(S_t) + \alpha \rho_{t:t+n-1} [G_{t:t+n} - V_{t+n-1}(S_t)]$$

, for $0 \leq t < T$

$\rho_{t:t+n-1}$ 就是 importance sampling ratio，是兩個policy使用這 n 個 action 的相對機率

$$\rho_{t:h} = \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$$

舉個例子，如果有任一個 action 不會被 policy π 採用 (i.e. $\pi(A_k|S_k) = 0$) 那這個 n-step return 會被忽略掉。在另一方面如果 這個如果被 π 挑中的機率大於 b ，那他的比重(weight)會大幅增加。若為on-policy 這兩個得值的比重永遠為1，所以我們先前的n-step Sarsa公式可以直接替換成：

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha \rho_{t:t+n-1} [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$

, for $0 \leq t < T$

Input: an arbitrary behavior policy b such that $b(a|s) > 0$, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be ϵ -greedy w.r.t Q , or as a fixed given policy

Algorithm parameter: step size $\alpha \in (0, 1]$, small $\epsilon > 0$, a positive integer n

All store and access operation (for S_t, A_t, R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq$ terminal

 Select and store an action $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal

 then $T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being update)

 If $\tau \geq 0$:

$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n, T-1)} \frac{\pi(A_i, S_i)}{b(A_i, S_i)} \quad (\rho_{\tau+1:\tau+n})$

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n}) \quad (G_{\tau:\tau+n})$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$

 If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is ϵ -greedy w.r.t Q

Until $\tau = T - 1$

7.5 Off-policy Learning Without Importance Sampling: The n-step Tree Backup Algorithm

這個點子是來自於右圖 3-step tree-backup backup。在主幹上有三個被標籤的 states 和 rewards 還有兩個 actions, S_t, A_t 是任意的初始 state-action pair。在旁邊的枝葉 state 是沒有被選擇到的 action, 因為我們沒有未選擇 action 的資料, 我們用 bootstrap 的方式使用殘枝的期望值。我們目前都會更新主幹的值, **而在 tree-backup 我們會額外去找主幹連接殘枝的值回來更新**, 這就是所謂的 tree-backup update。簡單來說我們會整合整棵樹的預測值再去做更新。

每個殘枝的 weight 會是在 policy π 的機率採取這個 action, 因此在第一層會貢獻 $\pi(a|S_{t+1})$ 除了實際上走的那條路 A_{t+1} 不會被計算在內, 而他的機率

$$\pi(A_{t+1}|S_{t+1})$$


, 會是第二層樹 action value 的比重。例如 a' 貢獻了

$$\pi(A_{t+1}|S_{t+1})\pi(a'|S_{t+2})$$

, 每個第三層樹則是

$$\pi(A_{t+1}|S_{t+1})\pi(A_{t+2}|S_{t+2})\pi(a'', S_{t+3})$$

以此類推。



我們把 3-step tree-backup update 拆解成6個half-step, 一半為 action 到接下來的 state, 另一半是考慮那個state的所有有可能的action 以及其在該 policy 發生的機率。

我們開始來推導 n-step tree-backup update 出這個演算法的細節。one-step return 跟 Expected Sarsa 一模一樣

$$G_{t:t+1} = R_{t+1} + \gamma \sum_a \pi(a|S_{t+1})Q_t(S_{t+1}, a) \quad (7.15)$$

若 $t < T - 1$ 那 2-step return 為

$$\begin{aligned} G_{t:t+2} &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1})Q_{t+1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1}) \left(R_{t+2} + \gamma \sum_a \pi(a|S_{t+2})Q_{t+1}(S_{t+2}, a) \right) \\ &= R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1})Q_{t+1}(S_{t+1}, a) + \gamma \pi(A_{t+1}, S_{t+1})G_{t+1:t+2} \end{aligned}$$

對於 $t < T - 2$ 後面一般化的公式如下：

$$G_{t:t+n} = R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1})Q_{t+n-1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1})G_{t+1:t+n} \quad (7.16)$$

對於在 $t < T - 1, n \geq 2$, $n = 1$ 已經包含在 (7.15) 了, 這個是用於 n-step Sarsa 對於一般 action value 的更新

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha[G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$

對於 $0 \leq t < T$ 當所有的 state-action pair 並未改變, $Q_{t+n}(s, a) = Q_{t+n-1}(s, a)$ 對於所有的 s, a 使得 $s \neq S_t$ or $a \neq A_t$ 。

n-step Tree Backup for estimating $Q \approx q_*$ or q_π

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be ϵ -greedy w.r.t Q , or to a fixed given policy

Algorithm parameter: step size $\alpha \in (0, 1]$, small $\epsilon > 0$, a positive integer n

All store and access operation (for S_t, A_t, R_t) can take their index mod $n + 1$

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$

Choose an action arbitrarily as a function of S_0 ; store A_0

$T \leftarrow \infty$

Loop for $t = 0, 1, 2, \dots$:

If $t < T$, then:

Take action A_t , Observe and store the next reward and state as R_{t+1}, S_{t+1}

If S_{t+1} is terminal

then $T \leftarrow t + 1$

else:

Choose an action arbitrarily A_{t+1} as a function of S_{t+1} ; Store A_{t+1}

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being update)

If $\tau \geq 0$:

If $t + 1 \geq T$

$G \leftarrow R_T$

else:

$G = R_{t+1} + \gamma \sum_{a \neq A_k} \pi(a|S_k) Q(S_k, a) + \gamma \pi(A_k|S_k) G$

Loop for $k = \min(t, T - 1)$ down through $\tau + 1$:

$G \leftarrow R_k + \gamma \sum_{a \neq A_k} \pi(a|S_k) Q(S_k, a) + \gamma \pi(A_k|S_k) G$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha[G - Q(S_\tau, A_\tau)]$

If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is ϵ -greedy w.r.t Q

Until $\tau = T - 1$
